Bhaskar Gurram
Zasti Inc.
Ashburn, VA 20147
USA

April 17, 2026

Dr. Jim Jansen
Editor-in-Chief
Information Processing & Management

Dear Dr. Jansen,

Thank you for the constructive feedback on our submission "AgentHallu-Bench: Measuring and Mitigating Error Propagation in Tool-Using Language Agents." We have made substantial revisions to address every point in your letter, and we are pleased to resubmit the revised manuscript, now titled **"Evaluating Tool-Using Language Agents: Judge Reliability, Propagation Cascades, and Runtime Mitigation in AgentProp-Bench."**

We renamed the benchmark from AgentHallu-Bench to AgentProp-Bench to reflect its focus on propagation measurement and to avoid confusion with the concurrently published "AgentHallu" paper (Liu et al., 2026, arXiv:2601.06818), which pursues step-level attribution on a different dataset.

Specific changes addressing your eight requests are detailed in the accompanying Response to the Editor and summarized here: (1) the abstract now leads with specific numbers rather than background (Abstract, ∼170 words); (2) explicit research objectives appear as a standalone Section 3; (3) a full dataset description is provided in Section 4 with domain breakdown, difficulty distribution, and example tasks; (4) a dedicated Discussion and Implications section (Section 7) addresses theoretical implications, practical guidance for agent designers, and relevance to IR/IS evaluation methodology; (5) the manuscript has undergone thorough copy editing for accessibility to non-native English speakers; (6) the paper has been expanded from approximately 5,200 to 9,500 words; (7) the literature review (Section 2) now includes 18 new 2024–2026 references from IR/IS and agent-hallucination venues, with at least three from SIGIR, CIKM, IP&M, and TOIS; (8) all references follow APA formatting via `elsarticle-num-names.bst`.

During revision, we also identified and corrected a methodological issue: the original stage-indicator definitions coupled two stages, producing a forced $r_{2,3} = 1.00$. We corrected the indicators, re-judged all traces with a validated 3-LLM ensemble, collected 100 stratified human labels ($\kappa = 0.432$ vs. ensemble), and re-ran the runtime interceptor under a concurrent no-interceptor control. All corrections are transparently disclosed in the Response to the Editor (Section 9) and in Section 5.2 of the revised manuscript.

We believe the revised manuscript is a strong fit for IP&M's audience, as the primary contribution—demonstrating that substring-based evaluation of agent outputs is at chance-level agreement with human annotation ($\kappa = 0.049$), while a validated LLM-ensemble judge reaches moderate agreement ($\kappa = 0.432$)—has direct implications for evaluation methodology in information retrieval and information science.

Per your request, we are willing to serve as reviewers for two manuscripts submitted to IP&M.

Best regards,

Bhaskar Gurram
Zasti Inc., Ashburn, VA, USA
gurrambhaskar.ai@gmail.com