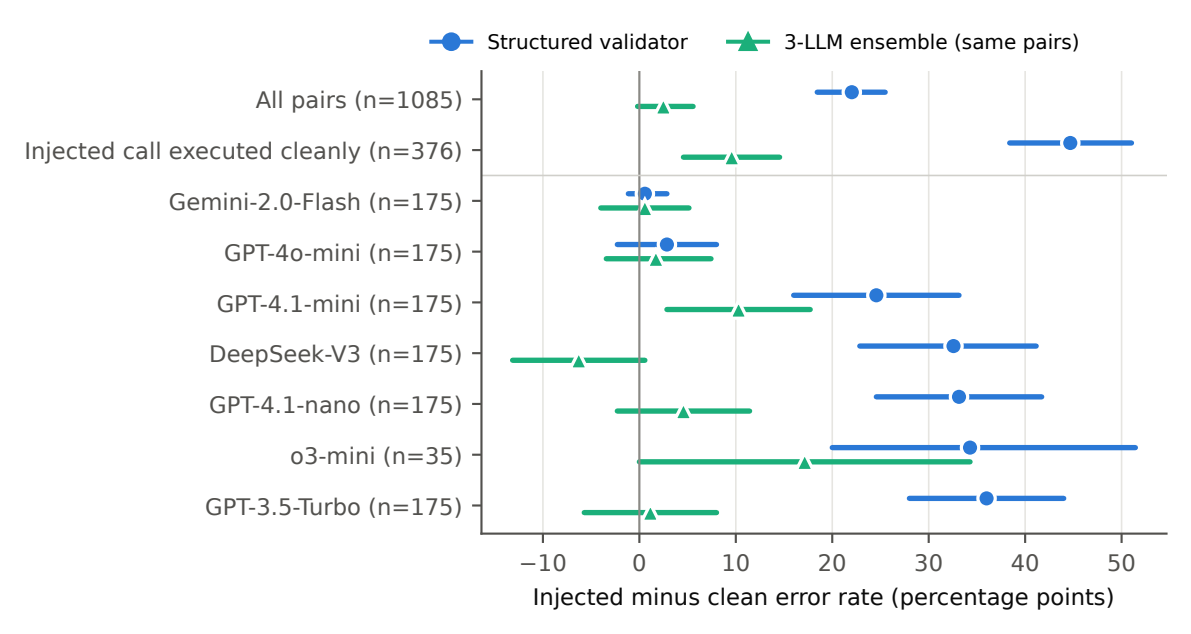
Structured validator
3-LLM ensemble (same pairs)
All pairs (n=1085)
Injected call executed cleanly (n=376)
Gemini-2.0-Flash (n=175)
GPT-4o-mini (n=175)
GPT-4.1-mini (n=175)
DeepSeek-V3 (n=175)
GPT-4.1-nano (n=175)
o3-mini (n=35)
GPT-3.5-Turbo (n=175)
−10
0
10
20
30
40
50
Injected minus clean error rate (percentage points)