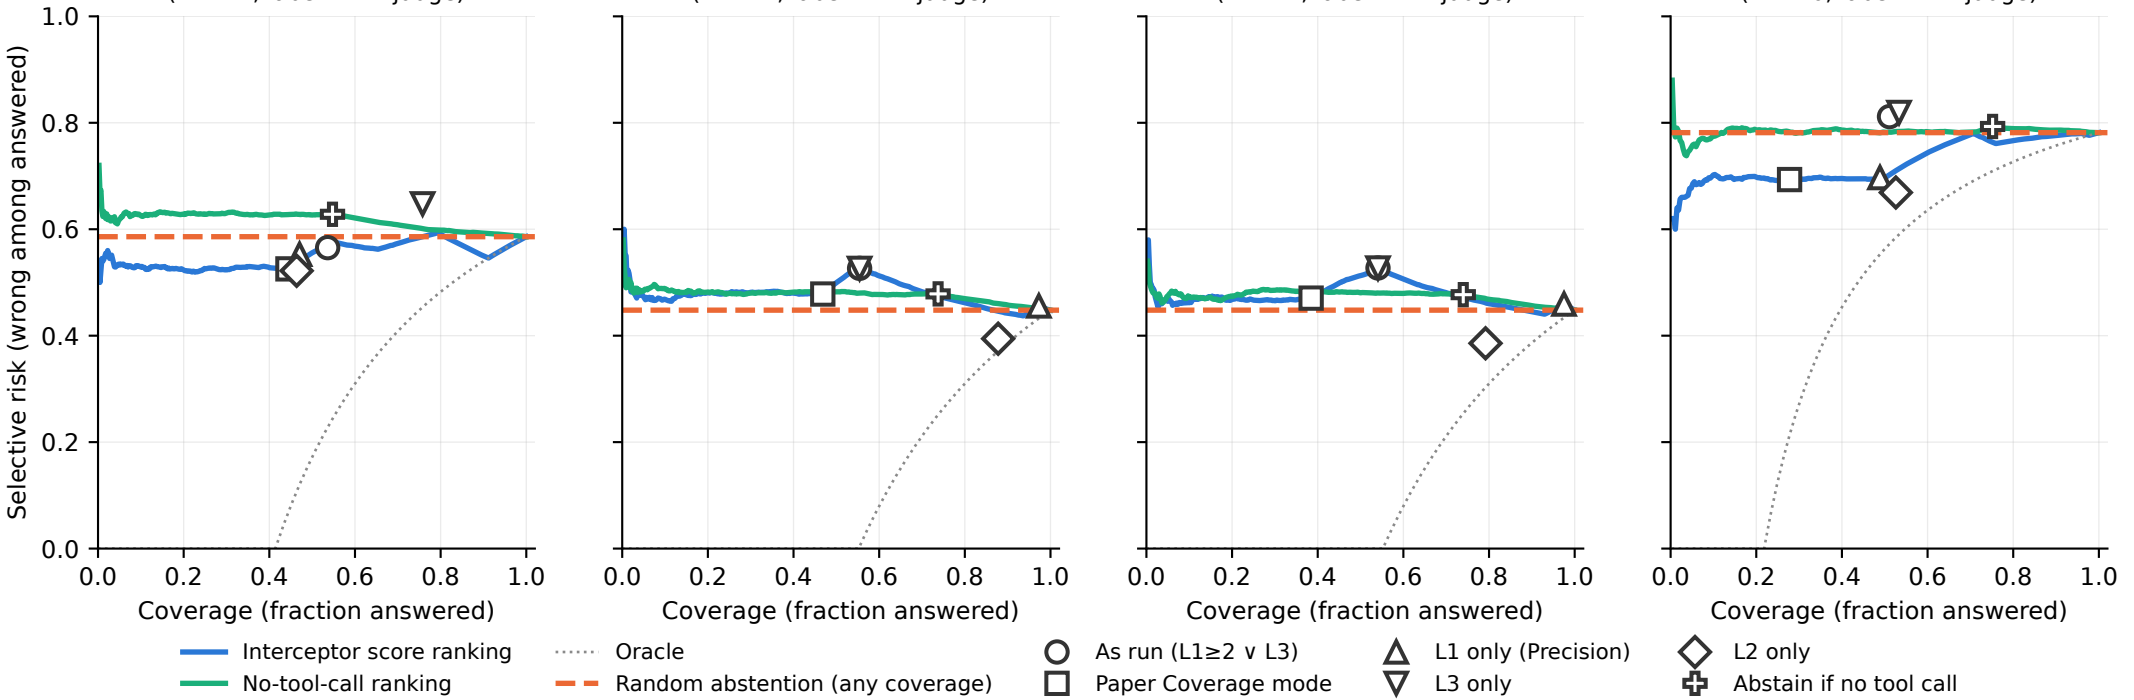
(a) Clean tasks, original harness
(n=442, label=LLM judge)
(b) Clean tasks, schema-exposed harness
(n=442, label=LLM judge)
(c) P2 semantic-wrong, schema-exposed
(n=442, label=LLM judge)
(d) Unseen retail domain, original harness
(n=270, label=LLM judge)
Selective risk (wrong among answered)
Coverage (fraction answered)
0.0
0.2
0.4
0.6
0.8
1.0
Interceptor score ranking
No-tool-call ranking
Oracle
Random abstention (any coverage)
As run (L1≥2 ∨ L3)
Paper Coverage mode
L1 only (Precision)
L3 only
L2 only
Abstain if no tool call